


 [torvalds](#) / [linux](#)

<> Code

 Pull requests 321 Actions Projects Security Insights v4.4-rc1 ▾

...

[linux](#) / [mm](#) / [madvise.c](#)

Mike Kravetz mm: madvise allow remove operation for hugetlbfs ...

 History 13 contributors

Raw

Blame

Copy File

30.46 KB 

551 lines (491 sloc) | 13.9 KB

```
1  /*
2  *      linux/mm/madvise.c
3  *
4  * Copyright (C) 1999  Linus Torvalds
5  * Copyright (C) 2002  Christoph Hellwig
6  */
7
8  #include <linux/mman.h>
9  #include <linux/pagemap.h>
10 #include <linux/syscalls.h>
11 #include <linux/mempolicy.h>
12 #include <linux/page-isolation.h>
13 #include <linux/hugetlb.h>
14 #include <linux/falloc.h>
15 #include <linux/sched.h>
16 #include <linux/ksm.h>
17 #include <linux/fs.h>
18 #include <linux/file.h>
19 #include <linux/blkdev.h>
20 #include <linux/backing-dev.h>
21 #include <linux/swap.h>
22 #include <linux/swapsops.h>
23
24 /*
25  * Any behaviour which results in changes to the vma->vm_flags needs to
26  * take mmap_sem for writing. Others, which simply traverse vmas, need
27  * to only take it for reading.
```

```
28  */
29  static int madvise_need_mmap_write(int behavior)
30  {
31      switch (behavior) {
32          case MADV_REMOVE:
33          case MADV_WILLNEED:
34          case MADV_DONTNEED:
35              return 0;
36          default:
37              /* be safe, default to 1. list exceptions explicitly */
38              return 1;
39      }
40  }
41
42  /*
43   * We can potentially split a vm area into separate
44   * areas, each area with its own behavior.
45   */
46  static long madvise_behavior(struct vm_area_struct *vma,
47                              struct vm_area_struct **prev,
48                              unsigned long start, unsigned long end, int behavior)
49  {
50      struct mm_struct *mm = vma->vm_mm;
51      int error = 0;
52      pgoff_t pgoff;
53      unsigned long new_flags = vma->vm_flags;
54
55      switch (behavior) {
56          case MADV_NORMAL:
57              new_flags = new_flags & ~VM_RAND_READ & ~VM_SEQ_READ;
58              break;
59          case MADV_SEQUENTIAL:
60              new_flags = (new_flags & ~VM_RAND_READ) | VM_SEQ_READ;
61              break;
62          case MADV_RANDOM:
63              new_flags = (new_flags & ~VM_SEQ_READ) | VM_RAND_READ;
64              break;
65          case MADV_DONTFORK:
66              new_flags |= VM_DONTCOPY;
67              break;
68          case MADV_DOFORK:
69              if (vma->vm_flags & VM_IO) {
70                  error = -EINVAL;
71                  goto out;
72              }
73              new_flags &= ~VM_DONTCOPY;
74              break;
75          case MADV_DONTDUMP:
```

```
76         new_flags |= VM_DONTDUMP;
77         break;
78     case MADV_DODUMP:
79         if (new_flags & VM_SPECIAL) {
80             error = -EINVAL;
81             goto out;
82         }
83         new_flags &= ~VM_DONTDUMP;
84         break;
85     case MADV_MERGEABLE:
86     case MADV_UNMERGEABLE:
87         error = ksm_madvise(vma, start, end, behavior, &new_flags);
88         if (error)
89             goto out;
90         break;
91     case MADV_HUGEPAGE:
92     case MADV_NOHUGEPAGE:
93         error = hugepage_madvise(vma, &new_flags, behavior);
94         if (error)
95             goto out;
96         break;
97     }
98
99     if (new_flags == vma->vm_flags) {
100         *prev = vma;
101         goto out;
102     }
103
104     pgoff = vma->vm_pgoff + ((start - vma->vm_start) >> PAGE_SHIFT);
105     *prev = vma_merge(mm, *prev, start, end, new_flags, vma->anon_vma,
106                     vma->vm_file, pgoff, vma_policy(vma),
107                     vma->vm_userfaultfd_ctx);
108     if (*prev) {
109         vma = *prev;
110         goto success;
111     }
112
113     *prev = vma;
114
115     if (start != vma->vm_start) {
116         error = split_vma(mm, vma, start, 1);
117         if (error)
118             goto out;
119     }
120
121     if (end != vma->vm_end) {
122         error = split_vma(mm, vma, end, 0);
123         if (error)
```

```
124         goto out;
125     }
126
127 success:
128     /*
129      * vm_flags is protected by the mmap_sem held in write mode.
130      */
131     vma->vm_flags = new_flags;
132
133 out:
134     if (error == -ENOMEM)
135         error = -EAGAIN;
136     return error;
137 }
138
139 #ifdef CONFIG_SWAP
140 static int swapin_walk_pmd_entry(pmd_t *pmd, unsigned long start,
141     unsigned long end, struct mm_walk *walk)
142 {
143     pte_t *orig_pte;
144     struct vm_area_struct *vma = walk->private;
145     unsigned long index;
146
147     if (pmd_none_or_trans_huge_or_clear_bad(pmd))
148         return 0;
149
150     for (index = start; index != end; index += PAGE_SIZE) {
151         pte_t pte;
152         swp_entry_t entry;
153         struct page *page;
154         spinlock_t *ptl;
155
156         orig_pte = pte_offset_map_lock(vma->vm_mm, pmd, start, &ptl);
157         pte = *(orig_pte + ((index - start) / PAGE_SIZE));
158         pte_unmap_unlock(orig_pte, ptl);
159
160         if (pte_present(pte) || pte_none(pte))
161             continue;
162         entry = pte_to_swp_entry(pte);
163         if (unlikely(non_swap_entry(entry)))
164             continue;
165
166         page = read_swap_cache_async(entry, GFP_HIGHUSER_MOVABLE,
167             vma, index);
168         if (page)
169             page_cache_release(page);
170     }
171 }
```

```
172         return 0;
173     }
174
175     static void force_swapin_readahead(struct vm_area_struct *vma,
176         unsigned long start, unsigned long end)
177     {
178         struct mm_walk walk = {
179             .mm = vma->vm_mm,
180             .pmd_entry = swapin_walk_pmd_entry,
181             .private = vma,
182         };
183
184         walk_page_range(start, end, &walk);
185
186         lru_add_drain();          /* Push any new pages onto the LRU now */
187     }
188
189     static void force_shm_swapin_readahead(struct vm_area_struct *vma,
190         unsigned long start, unsigned long end,
191         struct address_space *mapping)
192     {
193         pgoff_t index;
194         struct page *page;
195         swp_entry_t swap;
196
197         for (; start < end; start += PAGE_SIZE) {
198             index = ((start - vma->vm_start) >> PAGE_SHIFT) + vma->vm_pgoff;
199
200             page = find_get_entry(mapping, index);
201             if (!radix_tree_exceptional_entry(page)) {
202                 if (page)
203                     page_cache_release(page);
204                 continue;
205             }
206             swap = radix_to_swp_entry(page);
207             page = read_swap_cache_async(swap, GFP_HIGHUSER_MOVABLE,
208                                         NULL, 0);
209             if (page)
210                 page_cache_release(page);
211         }
212
213         lru_add_drain();          /* Push any new pages onto the LRU now */
214     }
215 #endif          /* CONFIG_SWAP */
216
217     /*
218     * Schedule all required I/O operations.  Do not wait for completion.
219     */
```

```

220 static long madvise_willneed(struct vm_area_struct *vma,
221                             struct vm_area_struct **prev,
222                             unsigned long start, unsigned long end)
223 {
224     struct file *file = vma->vm_file;
225
226 #ifdef CONFIG_SWAP
227     if (!file) {
228         *prev = vma;
229         force_swapin_readahead(vma, start, end);
230         return 0;
231     }
232
233     if (shmem_mapping(file->f_mapping)) {
234         *prev = vma;
235         force_shm_swapin_readahead(vma, start, end,
236                                   file->f_mapping);
237         return 0;
238     }
239 #else
240     if (!file)
241         return -EBADF;
242 #endif
243
244     if (IS_DAX(file_inode(file))) {
245         /* no bad return value, but ignore advice */
246         return 0;
247     }
248
249     *prev = vma;
250     start = ((start - vma->vm_start) >> PAGE_SHIFT) + vma->vm_pgoff;
251     if (end > vma->vm_end)
252         end = vma->vm_end;
253     end = ((end - vma->vm_start) >> PAGE_SHIFT) + vma->vm_pgoff;
254
255     force_page_cache_readahead(file->f_mapping, file, start, end - start);
256     return 0;
257 }
258
259 /*
260  * Application no longer needs these pages.  If the pages are dirty,
261  * it's OK to just throw them away.  The app will be more careful about
262  * data it wants to keep.  Be sure to free swap resources too.  The
263  * zap_page_range call sets things up for shrink_active_list to actually free
264  * these pages later if no one else has touched them in the meantime,
265  * although we could add these pages to a global reuse list for
266  * shrink_active_list to pick up before reclaiming other pages.
267  */

```

```
268 * NB: This interface discards data rather than pushes it out to swap,
269 * as some implementations do. This has performance implications for
270 * applications like large transactional databases which want to discard
271 * pages in anonymous maps after committing to backing store the data
272 * that was kept in them. There is no reason to write this data out to
273 * the swap area if the application is discarding it.
274 *
275 * An interface that causes the system to free clean pages and flush
276 * dirty pages is already available as msync(MS_INVALIDATE).
277 */
278 static long madvise_dontneed(struct vm_area_struct *vma,
279                             struct vm_area_struct **prev,
280                             unsigned long start, unsigned long end)
281 {
282     *prev = vma;
283     if (vma->vm_flags & (VM_LOCKED|VM_HUGETLB|VM_PFNMAP))
284         return -EINVAL;
285
286     zap_page_range(vma, start, end - start, NULL);
287     return 0;
288 }
289
290 /*
291 * Application wants to free up the pages and associated backing store.
292 * This is effectively punching a hole into the middle of a file.
293 */
294 static long madvise_remove(struct vm_area_struct *vma,
295                           struct vm_area_struct **prev,
296                           unsigned long start, unsigned long end)
297 {
298     loff_t offset;
299     int error;
300     struct file *f;
301
302     *prev = NULL; /* tell sys_madvise we drop mmap_sem */
303
304     if (vma->vm_flags & VM_LOCKED)
305         return -EINVAL;
306
307     f = vma->vm_file;
308
309     if (!f || !f->f_mapping || !f->f_mapping->host) {
310         return -EINVAL;
311     }
312
313     if ((vma->vm_flags & (VM_SHARED|VM_WRITE)) != (VM_SHARED|VM_WRITE))
314         return -EACCES;
315 }
```

```

316     offset = (loff_t)(start - vma->vm_start)
317             + ((loff_t)vma->vm_pgoff << PAGE_SHIFT);
318
319     /*
320      * Filesystem's fallocate may need to take i_mutex. We need to
321      * explicitly grab a reference because the vma (and hence the
322      * vma's reference to the file) can go away as soon as we drop
323      * mmap_sem.
324      */
325     get_file(f);
326     up_read(&current->mm->mmap_sem);
327     error = vfs_fallocate(f,
328                          FALLOC_FL_PUNCH_HOLE | FALLOC_FL_KEEP_SIZE,
329                          offset, end - start);
330     fput(f);
331     down_read(&current->mm->mmap_sem);
332     return error;
333 }
334
335 #ifdef CONFIG_MEMORY_FAILURE
336 /*
337  * Error injection support for memory error handling.
338  */
339 static int madvise_hwpoinson(int bhv, unsigned long start, unsigned long end)
340 {
341     struct page *p;
342     if (!capable(CAP_SYS_ADMIN))
343         return -EPERM;
344     for (; start < end; start += PAGE_SIZE <<
345          compound_order(compound_head(p))) {
346         int ret;
347
348         ret = get_user_pages_fast(start, 1, 0, &p);
349         if (ret != 1)
350             return ret;
351
352         if (PageHWPoison(p)) {
353             put_page(p);
354             continue;
355         }
356         if (bhv == MADV_SOFT_OFFLINE) {
357             pr_info("Soft offlining page %#lx at %#lx\n",
358                    page_to_pfn(p), start);
359             ret = soft_offline_page(p, MF_COUNT_INCREASED);
360             if (ret)
361                 return ret;
362             continue;
363         }

```



```
364         pr_info("Injecting memory failure for page %#lx at %#lx\n",
365                 page_to_pfn(p), start);
366         /* Ignore return value for now */
367         memory_failure(page_to_pfn(p), 0, MF_COUNT_INCREASED);
368     }
369     return 0;
370 }
371 #endif
372
373 static long
374 madvise_vma(struct vm_area_struct *vma, struct vm_area_struct **prev,
375             unsigned long start, unsigned long end, int behavior)
376 {
377     switch (behavior) {
378     case MADV_REMOVE:
379         return madvise_remove(vma, prev, start, end);
380     case MADV_WILLNEED:
381         return madvise_willneed(vma, prev, start, end);
382     case MADV_DONTNEED:
383         return madvise_dontneed(vma, prev, start, end);
384     default:
385         return madvise_behavior(vma, prev, start, end, behavior);
386     }
387 }
388
389 static bool
390 madvise_behavior_valid(int behavior)
391 {
392     switch (behavior) {
393     case MADV_DOFORK:
394     case MADV_DONTFORK:
395     case MADV_NORMAL:
396     case MADV_SEQUENTIAL:
397     case MADV_RANDOM:
398     case MADV_REMOVE:
399     case MADV_WILLNEED:
400     case MADV_DONTNEED:
401 #ifdef CONFIG_KSM
402     case MADV_MERGEABLE:
403     case MADV_UNMERGEABLE:
404 #endif
405 #ifdef CONFIG_TRANSPARENT_HUGEPAGE
406     case MADV_HUGEPAGE:
407     case MADV_NOHUGEPAGE:
408 #endif
409     case MADV_DONTDUMP:
410     case MADV_DODUMP:
411         return true;
```

```
412
413     default:
414         return false;
415     }
416 }
417
418 /*
419  * The madvise(2) system call.
420  *
421  * Applications can use madvise() to advise the kernel how it should
422  * handle paging I/O in this VM area. The idea is to help the kernel
423  * use appropriate read-ahead and caching techniques. The information
424  * provided is advisory only, and can be safely disregarded by the
425  * kernel without affecting the correct operation of the application.
426  *
427  * behavior values:
428  *  MADV_NORMAL - the default behavior is to read clusters. This
429  *                results in some read-ahead and read-behind.
430  *  MADV_RANDOM - the system should read the minimum amount of data
431  *                on any access, since it is unlikely that the appli-
432  *                cation will need more than what it asks for.
433  *  MADV_SEQUENTIAL - pages in the given range will probably be accessed
434  *                once, so they can be aggressively read ahead, and
435  *                can be freed soon after they are accessed.
436  *  MADV_WILLNEED - the application is notifying the system to read
437  *                some pages ahead.
438  *  MADV_DONTNEED - the application is finished with the given range,
439  *                so the kernel can free resources associated with it.
440  *  MADV_REMOVE - the application wants to free up the given range of
441  *                pages and associated backing store.
442  *  MADV_DONTFORK - omit this area from child's address space when forking:
443  *                typically, to avoid COWing pages pinned by get_user_pages().
444  *  MADV_DOFORK - cancel MADV_DONTFORK: no longer omit this area when forking.
445  *  MADV_MERGEABLE - the application recommends that KSM try to merge pages in
446  *                this area with pages of identical content from other such areas.
447  *  MADV_UNMERGEABLE - cancel MADV_MERGEABLE: no longer merge pages with others.
448  *
449  * return values:
450  *  zero - success
451  *  -EINVAL - start + len < 0, start is not page-aligned,
452  *            "behavior" is not a valid value, or application
453  *            is attempting to release locked or shared pages.
454  *  -ENOMEM - addresses in the specified range are not currently
455  *            mapped, or are outside the AS of the process.
456  *  -EIO - an I/O error occurred while paging in data.
457  *  -EBADF - map exists, but area maps something that isn't a file.
458  *  -EAGAIN - a kernel resource was temporarily unavailable.
459  */
```

```
460 SYSCALL_DEFINE3(madvise, unsigned long, start, size_t, len_in, int, behavior)
461 {
462     unsigned long end, tmp;
463     struct vm_area_struct *vma, *prev;
464     int unmapped_error = 0;
465     int error = -EINVAL;
466     int write;
467     size_t len;
468     struct blk_plug plug;
469
470     #ifdef CONFIG_MEMORY_FAILURE
471         if (behavior == MADV_HWPOISON || behavior == MADV_SOFT_OFFLINE)
472             return madvise_hwpoinson(behavior, start, start+len_in);
473     #endif
474     if (!madvise_behavior_valid(behavior))
475         return error;
476
477     if (start & ~PAGE_MASK)
478         return error;
479     len = (len_in + ~PAGE_MASK) & PAGE_MASK;
480
481     /* Check to see whether len was rounded up from small -ve to zero */
482     if (len_in && !len)
483         return error;
484
485     end = start + len;
486     if (end < start)
487         return error;
488
489     error = 0;
490     if (end == start)
491         return error;
492
493     write = madvise_need_mmap_write(behavior);
494     if (write)
495         down_write(&current->mm->mmap_sem);
496     else
497         down_read(&current->mm->mmap_sem);
498
499     /*
500      * If the interval [start,end) covers some unmapped address
501      * ranges, just ignore them, but return -ENOMEM at the end.
502      * - different from the way of handling in mlock etc.
503      */
504     vma = find_vma_prev(current->mm, start, &prev);
505     if (vma && start > vma->vm_start)
506         prev = vma;
507
```

```
508     blk_start_plug(&plug);
509     for (;;) {
510         /* Still start < end. */
511         error = -ENOMEM;
512         if (!vma)
513             goto out;
514
515         /* Here start < (end|vma->vm_end). */
516         if (start < vma->vm_start) {
517             unmapped_error = -ENOMEM;
518             start = vma->vm_start;
519             if (start >= end)
520                 goto out;
521         }
522
523         /* Here vma->vm_start <= start < (end|vma->vm_end) */
524         tmp = vma->vm_end;
525         if (end < tmp)
526             tmp = end;
527
528         /* Here vma->vm_start <= start < tmp <= (end|vma->vm_end). */
529         error = madvise_vma(vma, &prev, start, tmp, behavior);
530         if (error)
531             goto out;
532         start = tmp;
533         if (prev && start < prev->vm_end)
534             start = prev->vm_end;
535         error = unmapped_error;
536         if (start >= end)
537             goto out;
538         if (prev)
539             vma = prev->vm_next;
540         else /* madvise_remove dropped mmap_sem */
541             vma = find_vma(current->mm, start);
542     }
543 out:
544     blk_finish_plug(&plug);
545     if (write)
546         up_write(&current->mm->mmap_sem);
547     else
548         up_read(&current->mm->mmap_sem);
549
550     return error;
551 }
```